Report on the testbed results of the BattleMesh V10

Alessandro GNAGNI
Leonardo MACCARI
Gioacchino MAZZUCCO
Claudio PISA
Bastian BITTORF
Johan ??
??

Abstract

The Wireless Battle of the Mesh is an yearly event that brings together people from across the world to test and compare the performance of different routing protocols for ad-hoc and mesh networks, like Babel, B.A.T.M.A.N., BMX6, OLSR, 802.11s. Every year the community gathers and set-up a testbed on which the protocols are run, developed, debugged and tested, and some performance measures are extracted. While the initial spirit of the event was to set-up a competition between the protocols (as the name suggests), with time it changed into a moment of exchange of experience, collective development of innovations in the field of mesh networks and wireless open source networking software. This document reports on the experimental results of the tenth edition of the Battle of The Mesh event.

June 13, 2017

1   BattleMesh v10

The Battle of the Mesh (WBM), as the official website says:

> It is a tournament with a social character. If you are a mesh networking enthusiast, community networking activist, or have an interest in mesh networks you might want to check this out!

> The goal of the WirelessBattleMesh events is to set-up hands-on testbed for each available mesh routing protocol with a standard test procedure for the different mesh networks. During the different WBM events, similar hardware and software configuration will be used based on the OpenWRT BoardSupportPackage and packages for each protocol implementation. The WBM events are also a great opportunity to develop testing tools for PHY/MAC radio layers (drivers, scripts and PHY analyzers).

WBM is now at its tenth edition, and is organized by a motivated and large group of people (approximately 80-100 participants in the whole week in the last editions from 2-3 continents).

2   The Testbed

This year, the testbed was realized with 17 TP-Link WDR4300 routers, equipped with two wireless interfaces (operating in the 2.4 and 5.0 GHz bands), and 5 Ubiquiti Unifi AC Pro. The latter were used to perform local testing and support to the tests, but did not participate to the routing, since they are equipped with a different chipset. The nodes were configured to set-up two independent networks, a "management" network, running on the 2.4GHz, and a "testing" network, running on the 5 GHz. The management network was configured with the IEEE 802.11s protocol, and was used only to access the nodes and perform tasks. The testing network was made with interfaces configured in ad-hoc mode and, for each test, was using a specific routing protocol.

Figure 1 reports the topology of the network as exported by one of the network node, when running the OLSRv1 protocol. This topology was used to understand and roughly guess the property of the network. The underlying image is the map of *Volkskundemuseum Wien*[1] that hosted the event, roughly, the distance from node 9 to node 17 is about 75m, and the width of the main room is about 8m (the main room is the room where nodes 1,4,7,6,31 are placed and where the conference took place).

Routers numbered from 1 to 17 are the TP-Link, while router 31 is an Ubiquiti router that is used in this case to extract the network topology. In the real tests this router does not participate to the network functioning. The transparency of the links represents the badness of the link quality as reported by the OLSRv1 protocol, using the ETX metric. A solid link means ETX = 1 (good link), while a transparent link means ETX > 1 (the highest, the worse).

Note that the topology in fig. 1 does not necessarily represent the topology used by other routing protocols, but gives an approximate idea of what links are directly

---

[1] The Austrian Museum of Folk Life And Folk Art, Laudongasse 15-19, Vienna, Austria
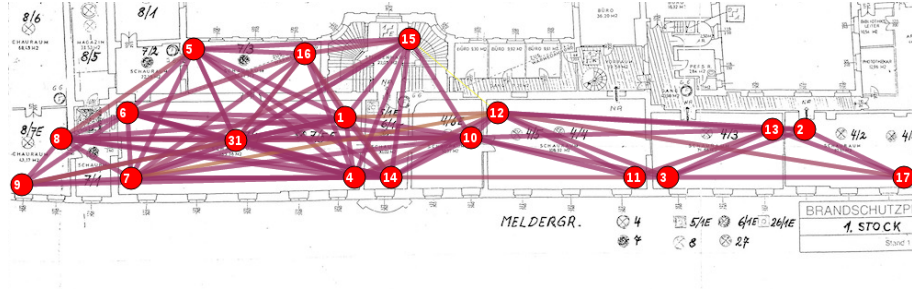
Fig. 1: Network topology, as exported from the OLSRv1 protocol

connected. Based on this topology, exported in the netJSON format[2], we calculated the weighted shortest path between any couple of node, using the networkX python graph library. We repeated the process twice, with two different transmission power level for the testing network. With the transmission power set to 17dBm the network was considered too dense, and thus of little interest for the routing function. Therefore the power was lowered to 10dBm, and produced the distribution of path costs reported in fig. 2. Each entry is computed as follows, given the network graph and a couple of nodes (A, B), the shortest path between A and B is computed via networkX and then the sum of the ETX values for the path is done[3]. We selected four nodes to perform the tests, node 8, 9, 17, and 2, that are at the extreme ends of the topology. fig. 2 reports the corresponding shortest path costs in the ranking.

2.1   The Protocols and the Experiments

Several protocols were tested during the WBM, not all the protocols (or their variants) have been tested on all the configuration. Table 1 reports the list of the protocols, a brief description, and the link to the source code.

Four out of six days that make the WBM were devoted to set-up the testbed, and only the last two were dedicated to the testing itself. The procedure of set-up and testing is error prone due to a vast set of reasons that range from the need to use a recent OpenWRT/LEDE version, the potential incompatibility with the last version of the protocols, their configuration, and last, but not least, the eventual bugs that are found while testing and need to be fixed. This is a very important part of the WBM, possibly the most important under the technical point of view (the social side of the WBM as equally important). During the tests the developers of the protocols (that are generally present at the event) test new features, compare different strategies and inevitably stumble upon unknown bugs in their protocols. This process is vital for them to improve their software and stabilize their code, and is arguably even more important than the results of the tests themselves.

---

[2] A generic format for exporting a network topology, see http://netjson.org
[3] For the code used for this task, see https://github.com/battlemesh/battlemeshv10-testbed/blob/master/tests/failure_recovery/parser_scripts/parse_json.py

Fig. 2: The ranked list of the ETX weight of all the shortest paths in the network.

Four set of experiments were planned, and three were fully performed:

- Ping test: a session of 100 pings from node 8 to 17, 8 to 2, 9 to 17, and 9 to 2 were performed to measure the loss and the delay distribution. This test was repeated only once.

- iperf test: a batch of 10 TCP iperf sessions were run (10s each) between node 8 and 17 and from node 9 to node 2. This experiment was repeated with two more variants, one in which node 14 and 10 were running another iperf with a 5Mbps limit, and a second one in which 4 5Mbps session was running between other nodes. These added sessions were not recorded but where used to increase the level of congestion in the central part of the network.

- Airtime Fairness tests: in this case we repeated the iperf tests with and without the Airtime Fairness enabled.

- pingall test: in this experiment, a random subset of all the possible couples was taken and an iperf session was performed. This test should have been repeated for all the protocols, with and without Airtime fairness enabled, but was not possible to complete.

Name	Desciption
BABEL	Distance-vector LIII routing protocol
BATMAN Advanced v4	Source Routed LII routing protocol
BATMAN Advanced v5	Source Routed LII routing protocol (experimental)
BMX7	Source Routed LIII routing protocol
BMX7TUN	BMX7 with IP tunnel support
OLSRv1	Link state LIII routing protocol
OLSRv2	Link state LIII routing protocol
OLSRv2_MPR	OLSRv2 with MPR enabled

Tab. 1: The list of tested protocols

3 The Results

3.1 Ping Tests

Figure 3 reports the percentage of lost packets in the ping tests. It shows that OLSRv1 and BMX are the ones that choose paths that are more conservative, so they deliver all or almost all their pings. OLSRv2, BATMAN4 and, to a lower extent BABEL are the ones that, instead, tend to choose paths that are more lossy. BATMAN5 largely underperforms compared to the others. This is an example of a typical situation in the WBM. BATMAN5 does not have a stable release, and the developers brought to the WBM a testing version, to initially study its performance. During the tests, these outlier performances made it possible to spot the presence of previously unknown software bugs, which become visible only when tested on networks larger than a certain size. It was not possible to patch BATMAN5 before the end of the WBM and thus, from now on the results of BATMAN5 are omitted.

Figure 4 reports data about the distribution of the RTT measured with ping. What clearly emerges is that BMX is consistently using paths with larger delays, and that BABEL on three cases over 4 has a very large range of values. OLSRv2 performs worse than OLSRv1 in the majority of the cases, with delays comparable to BATMAN. Note that OLSRv2, BATMAN, and BABEL are advantaged in this comparison, since they have non-zero loss. It is reasonable to assume that the packets that could not be delivered, if they could reach the destination would probably have a large RTT.

Finally, figs. 5 to 8 report the whole set of RTT measured for each run. Note that the tests were done in sequence, so the network conditions may have varied from one test to the next one, note also the log scale on y axis. It is clear that there are two phenomenons, one is the difference from one protocol to another that impacts the median value, the other is the presence of many outliers that influence the whiskers of fig. 4.

Fig. 3: The percentage of lost packets per protocol

Fig. 4: The values of the 10th percentile, 90th percentile, and median value of the RTT for all the pings

Fig. 5: The values of the RTT per each ping, 9 to 17


Fig. 6: The values of the RTT per each ping, 9 to 2

Fig. 7: The values of the RTT per each ping, 8 to 17

Fig. 8: The values of the RTT per each ping, 8 to 2

3.2 iperf Test

Figures 9 to 11 report the performance of all the protocols when two parallel iperf sessions are run to measure the performance and zero, one or 4 sessions are added to increase background noise. The background sessions are fixed at 5Mbps and take place between couples (14,10), (4,15), (1,16), (12,11). Some observations that can be done are:

- There is a net performance degradation from one figure to the next one.

- With zero and one flow, in the path 8-17 (that according to fig. 2 is the shortest of the two) the protocol behave similarly, while in the path 9-2 there is a higher performance variation. In particular, BMX7 consistently performs better than the others (as the median value) but also shows the largest deviation. This may be the consequence of the combination of the zero loss and the stable delay that BMX7 is able to obtain (see figs. 3 and 4).

- with 5 flows, the relative difference increases and OLSRv2 seems to perform better than the other protocols averaging the results of both paths.

Fig. 9: iperf tests without background traffic.

Figure 12 reports the comparison of the performance of the iperf sessions, without background traffic with and without the Airtime fairness enabled. Each color in the figure corresponds to a protocol, the left value (with a dot at the median value) is the performance without the Airtime fairness enabled, the right value (with a dash at the median value) is the performance with the Airtime fairness enabled. It is hard to draw

Fig. 10: iperf tests with one background traffic flow.

Fig. 11: iperf tests with four background traffic flows.

any conclusion on the importance of Airtime fairness, BMX7 still remains the protocol that guarantees the highest throughput, even if it is affected by the Airtime fairness in the opposite way in the two paths. Also the other protocols are affected both positively and negatively by Fairness, so it is unclear if the results are the effect of fairness, or of the changed conditions in the testbed from one run to the other. The only conclusion is that Airtime fairness does not influence the performance of mesh protocols in an critical way.


Fig. 12: iperf tests without background traffic and with/without airtime fairness.

4 Suggestion for the Future Battlemeshers

What follows is a set of recommendations we, as the people that spent most time on the testbed, want to give to the ones that in the next editions of the battlemesh will set-up the testbed.

Assign micro-tasks to other people

Setting up the testbed is full-time activity, so it is demanding on those that do it. Still, some tasks can be "outsourced" to others, tasks that are small, atomic and tedious, such as writing scripts to parse the output of commands used for testing (ping, iperf, traceroute...). When needed, instead of spending tens of minutes to write the script, it is useful to ask for help to others, this was done a couple of times and worked good. Ideally: write the task to an etherpad, ring a bell, and wait for someone to provide a script that does the job.

Use cables!

We choose 802.11s for the management network because, if it does not work, we can't blame any of the protocols under test. 802.11s works decently but, as any other solution is influenced by the state of the network. We believe that the best thing would be to buy a 300M reel of ethernet cable, and use that to manage the nodes. This would spare us of many failures in connecting, launching scripts, failing, restarting etc.... Powerline could be an alternative.

Do small tests

Large scale tests (all nodes against all nodes) are tempting, but since the network becomes complex it is hard to understand what is happening. We decided to start with small tests, where it is easier to understand what happens, while it happens.

Bring hardware for night-time tests

Probably the best moment to perform large-scale tests is during the night. But no one wants to leave his/her laptop at the BM nighttime. So it is important to remember to bring some other hardware (like a few raspberries) that can be used to perform automated tests during the night.

Fork the testbed

It happens that some of the devels need to use the testbed to fix bugs that happen during the experiments, and it is the added value of BM for them. On the other side, we can not freeze testing for long while they debug/improve their code. So a wise idea is to realize a small tbed (4-5 nodes) detached from the large one that devels can use to fix their bugs.

5   Conclusions

While there is a discussion in the BM community about changing the format of the event to something that is more conference-wise and less testbed-wise, our impression is that the testbed is important. Even if the results are incomplete, and scientifically not always sound, the whole process of setting up the testbed, running the protocols, debugging with the developers is an integral part of the conference and guarantees that the experts will participate to the conference, which makes it different from any other conference.